

When Should AI Assist the Assessor?

A Philippine Higher-Education Benchmark of Failure Containment, Control
Allocation, and Human Verification Burden

Bea Charmelyn T. Lambitco, FRM
Philippines · AI for Asia Fellowship 2026
September 18, 2026

Executive Summary

Generative artificial intelligence (AI) can help faculty and staff in Philippine higher education interpret how a grading, attendance, completion or academic-integrity rule applies to a particular case. Its usefulness depends on more than a plausible answer: the answer must apply the correct institutional rule, stay within the available evidence and leave consequential decisions with the authorized person; checking it also takes time. This study asks when AI should assist the people who apply assessment rules, and what they must still verify themselves.

The study built 24 test cases, each with an answer key, from the published rules of nine Philippine higher-education institutions (HEIs). Two AI model configurations answered every case twice, producing 96 answers, and each answer was released in two ways: directly, and after four software checks that could block it or send it to a person. Because the same answer appears in both routes, any difference in outcomes is attributable to the checks. Independent reviewers and an AI judge cross-checked the scoring, and checking time was recorded for 24 answers.

Of the 96 answers, 60 (15 complete cases) were scored, among them all 16 that the checks withheld. All 60 were acceptable and none contained a serious error, so the checks had no error to stop and every answer they held back was sound: eleven were routed because the model labeled its own acceptable answer as a grade or approval decision, two were blocked over an unrecognized unit label on a correct number, and three for quoting the syllabus under review. Of 419 quotations offered across all 96 answers, every one was present in the material supplied to the model. Median decision times ranged from 98 to 158 seconds across conditions, and 7 of 20 timed answers needed some correction. Three independent reviewers and an AI judge agreed with every acceptability decision they could compare. In an exploratory extension, the checks released two confirmed calculation errors from one Claude model because the model reported the number under a different name.

On these cases, with the governing rules supplied, the answers held up and the checks were the weak link: they reacted to labels rather than to errors, and every answer they withheld was one a user could have acted on. Institutions should also ensure that their grading and integrity rules are dated and versioned, record which office decides exceptions, and keep final decisions with accountable people; the Commission on Higher Education (CHED) and accreditors can ask for evidence of versioned rules and decision authority. Member states of the Association of Southeast Asian Nations (ASEAN) can reuse the method with cases rebuilt from their own rules and languages. The findings are limited to these cases, configurations and prompt design, a single researcher who also developed the cases, and partial scoring; they do not show that AI can assess students on its own or save teachers time.

1. Introduction

Grading is only one part of assessment. Around each grade sit smaller decisions: whether a syllabus states an adequate AI-use policy, whether a student has exceeded an absence limit, whether an incomplete grade has lapsed, and whether a request needs a dean's approval. Each decision combines a written rule, some interpretation and a question of who has the authority to decide. An artificial intelligence (AI) assistant can draft a recommendation from a handbook excerpt in seconds, but a correctly computed percentage may still carry an unsupported consequence.

Research on generative AI in higher-education assessment has examined academic integrity, grading and feedback (Jukiewicz & Wyrwa, 2026; Xia et al., 2024), and a scoping review advised institutions to revisit their assessment policies (Xia et al., 2024). The literature reviewed for this study says much less about the next step, when a written policy is applied to an individual case and the result must be checked before anyone acts on it. Advice to keep a human in the loop seldom specifies what that person should check, and verification effort should be reported alongside correctness (Crescitelli et al., 2026). The gap matters most where staff have little time and technical support, because a fast draft can still leave them to find the source, confirm the interpretation and correct the errors; this study therefore treats checking and correction as part of the cost of assistance.

The study has three objectives: to compare how two AI model configurations answer rule-based assessment requests and how simple software checks change which of the same answers are released; to identify which errors those checks catch, which they miss and which decisions must remain with an accountable person; and to describe the checking and correction work that remains and its implications for institutions with limited resources. [Appendix B](#) maps each of the fellowship's guide questions to the section that answers it.

2. AI Landscape and Research Context

2.1 Philippine guidance and the problem of applying it

Philippine higher education operates within a developing AI-governance landscape, in which reported use has moved faster than guidance. The National AI Strategy Roadmap 2.0, launched in July 2024, builds on an information technology and business process management sector of more than 1,300 companies and a young, English-speaking workforce with over 800,000 college graduates each year, and it calls for AI governance that can adapt as the technology changes (OECD.AI, 2026). Readiness is uneven: a 2025 assessment by the United Nations Educational, Scientific and Cultural Organization (UNESCO) found that no dedicated lead agency consolidated national AI policy and identified the lack of affordable broadband as a root constraint on building digital skills (UNESCO, 2025). According to a news report on initial findings from an Association of Southeast Asian Nations (ASEAN) Foundation study, 83.4% of Filipino students and 73.1% of educators said they used generative AI tools in education, although the report does not describe the sample or method (Untalan, 2026).

For higher education, a 2026 discussion paper from the University of the Philippines Center for Integrative and Development Studies found no national rules yet covering AI use in all higher-education institutions (HEIs), with the first concrete guidelines written at the University of the Philippines and its Open University (Gervacio, 2026). Toquero's preliminary analysis of public documents, websites and statements describes institutional AI policies as fragmented and often left to individual departments (Toquero, 2026). A research paper prepared for the Second Congressional Commission on Education (EDCOM II) recommends that HEIs formulate and publish AI-use policies supported by templates from the Commission on Higher Education (CHED) (Vergara, 2024), as Filipino instructors and administrators have also asked (Giray et al., 2024). At CHED's national convention on AI in higher education in February 2026, according to a news report, the CHED chair said that adopting AI must not compromise academic integrity, and a commissioner said that CHED wants HEIs to write internal AI policies (Oladive, 2026); the convention also discussed a proposed national AI framework (Commission on Higher Education Regional Office XI, 2026).

Even where policies exist, students may not know them. In a voluntary survey of 451 students at a rural state university in Cebu, all respondents had used AI tools, but only about 17% knew their institution's AI-use policy (Villarino, 2025). A search conducted for this study across 22 institutions and campuses found AI-specific guidance at some, only general or undated handbooks at others, and nothing usable at a few; this reflects what could be found online, not which policies exist. A search of the titles of CHED memorandum orders issued from 2023 to 2026, as listed on the accessible CHED web pages, found none on AI; memorandum circulars and advisories were not systematically searched ([Appendix I](#)).

Rules written for other settings show the safeguards that regulators expect. The registry annex of the 2026 Department of Education (DepEd) guidelines for basic education classifies AI used for grading as high risk and allows it only with safeguards, including a privacy impact assessment, human oversight, audit logs and a means of appeal (Department of Education, 2026). The Philippine Accrediting Association of Schools, Colleges and Universities (PAASCU), a voluntary accreditor, requires human experts to make final accreditation decisions and requires disclosure of AI use in accreditation (Philippine Accrediting Association of Schools, Colleges and Universities, 2025). Neither instrument governs the college cases in this study, which apply each institution's own published grading, completion and integrity rules.

2.2 Who benefits and who is left behind

The benefits are unlikely to be shared evenly, and internet access illustrates why. According to a news report on the 2024 National Information and Communications Technology (ICT) Household Survey, 48.8% of households had internet access at home, ranging from 68.7% in the National Capital Region to 21.2% in Zamboanga Peninsula, with 27.7% in Bangsamoro (SunStar Cebu, 2025). Institutions that can keep their rules current and maintain software checks may therefore gain more, although this study did not measure that effect.

Errors can also fall unevenly on students. Two cases in this study ([Appendix A](#)) involve students with little room for error: a scholarship grantee facing a shorter completion window (W2-08) and a student who missed a final examination because of illness (W1-07). In a 2023 study, seven AI-text detectors wrongly flagged, on average, more than half of 91 essays by non-native English writers as AI-generated (Liang et al., 2023), a concern wherever students write in English as an additional language; case W2-03 applies a published university rule that does not allow a detector result to be the sole basis for a finding of dishonesty. Where automated decisions by AI systems pose significant risks to data subjects, a National Privacy Commission advisory requires personal information controllers, such as HEIs processing student records, to provide meaningful human intervention and a means to contest those decisions (National Privacy Commission, 2024).

2.3 What research establishes and leaves open

Frameworks help educators decide what AI use is permitted. The Artificial Intelligence Assessment Scale specifies the level of AI use permitted in a task (Perkins et al., 2024), but Corbin et al. (2025) argue that such permission statements cannot enforce themselves, and a global Delphi study lists human oversight and accountability among the core areas of higher-education AI policy (Crompton et al., 2026). Permission to use a tool and evidence that its recommendation is dependable are, however, separate questions. Evidence on AI grading is mixed: a review reports stronger results on bounded tasks and judges hybrid, teacher-verified grading most effective (Jukiewicz & Wyrwa, 2026), and in a simulation, sending only uncertain automated grades to a person improved accuracy (Korthals et al., 2026).

Human review has limits of its own. AI explanations improve human decisions mainly when they allow a person to verify an output (Fok & Weld, 2024); people rely more on AI when verification is costly (Vasconcelos et al., 2023); and designs that force deliberate engagement reduce overreliance but are less liked (Bućinca et al., 2021). Teachers in Australia and Sweden also described unseen time spent checking and fixing AI outputs (Selwyn et al., 2025).

Supplying the rules does not ensure that a model applies them: in a rule-guided reasoning benchmark, models missed applicable rules, confused similar ones and made arithmetic errors (Zhou et al., 2025). Software can catch some of these errors, for example by recomputing a weighted average or confirming that a quotation exists; others require interpretation, or recognizing that a decision belongs to an authorized person, the classic distinction between automating analysis and automating decisions (Parasuraman et al., 2000). Models also do not reliably hold back: in a benchmark of 20 frontier models, they often answered questions they should have declined (Kirichenko et al., 2025), and research on abstention weighs fewer wrong answers against refusing too often (Geifman & El-Yaniv, 2019; Wen et al., 2025). Withholding has costs as well; in a clinical experiment reported in a preprint, withholding AI predictions judged unreliable reduced harm but increased missed diagnoses (Jabbour et al., 2026). Useful answers that are withheld must therefore be counted alongside errors that are stopped.

Outside education, guard code derived from a written policy outperformed prompt reminders on 22 policy-violating tasks in a simulated airline benchmark (Zwerdling et al., 2025), and hosted models vary across runs even with fixed settings (Atil et al., 2025). In a preprint benchmark, correctness judgments were only as reliable as the reference answers, and AI-generated references were often wrong (Krumdick et al., 2026); this is why the present study reports independent reviewer and AI-judge agreement alongside its own scores. A Philippine prototype already answers student-

handbook questions with a language model (Ilagan & Ilagan, 2024), but the literature search ([Appendix I](#)) found no Philippine study that tests software checks on the same AI answers, draws decision authority from named institutions' rules and measures the checking work that remains.

3. Research Approach

3.1 Research questions and paired design

The research questions, fixed before any answers were collected, are reproduced verbatim below. In their wording, a large language model (LLM) configuration means a model used through a named interface and setting, and deterministic release controls means rule-based software checks that always give the same result for the same answer (hereafter, the checks). The main research question (RQ) is: **For selected rule-governed assessment-support workflows in Philippine higher education, how do frontier and efficient LLM configurations, with and without deterministic release controls, compare in correctness, failure containment, inappropriate delegation, useful release, and associated model cost and single-researcher handling burden?**

- **RQ1.** How do the selected frontier and efficient models differ in correctness and serious failures, and how does applying deterministic release controls to the same responses affect serious-failure release, useful release, and unnecessary withholding?
- **RQ2.** Which failure classes can deterministic validators detect and contain, where do those validators fall short, and what roles for language-model interpretation and accountable human review are supported by the benchmark findings and governing sources?
- **RQ3.** What model cost and single-researcher verification and remediation burden are associated with each configuration, and what preliminary implications do these findings have for resource-constrained assessment-support workflows in Philippine higher education?

Failure containment means that serious answers are withheld from ordinary release; control allocation means deciding which judgments go to software checks, to model interpretation or to a person; and human verification burden means the recorded time to decide on and correct an answer. Each configuration answered each case twice, and each of the resulting 96 answers was released both directly and under the checks, giving 192 paired records rather than 192 independent responses. Because the checks never rewrite an answer, any difference between the two routes is attributable to the checks and not to a better second attempt. To test whether the checks add value, they are also compared with three simple baseline policies applied to the same saved answers: withhold every answer; block only answers with format errors; and send every answer in the 19 cases that call for judgment to a person. A policy is considered better than another only if it releases no more serious answers and no fewer useful ones.

3.2 Cases and answer keys

The 24 cases fall into three workflow groups of eight: evidence, source and version (W1); delegation and referral, that is, whether the assistant leaves a decision to the authorized person (W2); and calculation and configuration (W3). Together they carry 90 scoring criteria ([Appendix A](#)). The nine institutions come from the National Capital Region, Central Visayas, Davao and Caraga and comprise one national university, two state universities, one local college and five private HEIs; Luzon outside Metro Manila, the other Visayas regions and most of Mindanao, including Bangsamoro, are not represented, and some cases share sources.

Each answer key contains a reference response, acceptable variants and scoring criteria linked to retained source passages. The scenarios and keys were drafted with AI assistance, revised by the researcher against the primary sources, challenged by a supplementary AI audit and re-reviewed after a delay ([Appendix D, Figure D1](#)). Astra, an evaluated configuration, helped refine the cases, and the planned review by independent Philippine higher-education experts did not occur before the freeze, so the keys rest on the researcher's judgment and disagreement with a key cannot simply be read as model error. All study materials were frozen on September 15, 2026, that is, locked and recorded with file hashes so that any later change would be detectable.

3.3 Generation and software checks

The configurations were recorded as Astra and Luna in OpenAI’s Codex interface at its Medium reasoning-effort setting. The intended models were gpt-6-astra and gpt-5.6-luna, but backend identifiers could not be verified, so the labels “frontier” and “efficient” describe the planned comparison rather than measured capability. Answers were collected manually, each in a fresh task; because tool and workspace access were not recorded, access to project files, including answer keys, cannot be excluded. Each prompt supplied the scenario, the task, the policy excerpt and a required JavaScript Object Notation (JSON) output format, and excluded the answer key. Required output formats can reduce reasoning quality (Tam et al., 2024), but any such effect is identical in both routes.

Table 1 lists the four frozen checks. Their parameters were taken from the source or the scenario, never from the answer key, and each check is deliberately narrow.

Table 1. *The four software checks*

Check	What it verifies	What it cannot see	Cases
Format	Valid JSON with the required fields, case ID and an allowed action	Whether the content is correct	24
Quotation	Each quoted passage appears in the supplied excerpt	Whether the passage supports the advice	24
Arithmetic	Named quantities, recomputed from scenario and source values	Renamed or omitted quantities	5
Routing	A consequential action the model declares (grade, approval, penalty)	Consequential advice given only in prose	24

The checks lead to one of three decisions. An answer that fails any check is blocked, and blocking takes precedence over routing. An answer that passes every check but declares a consequential action is routed to a person, who decides. All other answers are released with a standard caution, which counts as release. A matched quotation shows only that the text exists in the source, not that it supports the recommendation; studies of AI legal research tools document the same gap (Magesh et al., 2025). Routing relies on the model’s own action label, and exact matching can block answers that differ harmlessly, for example in unit labels or rounding.

3.4 Outcomes and analysis

Every scored answer received criterion-level judgments and a severity rating before its release decisions were matched to the scores. Table 2 defines the severity levels and the four release outcomes derived from them.

Table 2. *Severity levels and release outcomes*

Term	Definition
Serious error	Critical: recommends an unsupported consequential action, such as a penalty or grade change (inappropriate delegation); Major: changes the substance or prevents a usable answer. Minor errors do not change the substance
Acceptable answer	No serious error and no undecided criterion; correct advice to obtain missing evidence can be acceptable
Useful release	An acceptable answer that reaches the user

Term	Definition
Unnecessary withholding	An acceptable answer that is routed or blocked (reported separately)
Serious released failure; serious answer withheld	A serious answer that reaches the user; one that is routed or blocked (no later correction assumed)

Unscored answers remain in the denominator of 48 per condition, and because they can count neither as serious nor as useful, coverage and best- and worst-case bounds are reported in [Appendix G](#). Following benchmark reporting guidance (Bean et al., 2025; Kapoor et al., 2024; Reuel et al., 2024), the analysis reports counts with their denominators and compares the two routes case by case, treating repetitions as related (Miller, 2024); case-resampled intervals can be too narrow for small benchmarks and are therefore labeled approximate (Bowyer et al., 2025). Failure classes were not coded during scoring, a departure from RQ2’s wording, so RQ2 is answered by comparing each check with the answers it withheld. No success threshold or significance test was set. Before any score was locked, an audit of the frozen analysis code found reporting errors; the corrections were recorded in a dated addendum, published in the project repository, and applied at the score lock on September 18, 2026.

3.5 Scoring, checking time and ethics

Answers were scored with model identity hidden, in a balanced order fixed in advance by random seed. At a pre-set cutoff on September 18, 2026, 48 answers covering 12 complete cases were scored and locked; a second stage, ordered and published before it began, then scored every answer the checks had withheld and no one had scored. Scoring stopped on September 22, 2026 at 60 answers in 15 complete cases, covering all three workflow groups and all nine institutions, none marked as impossible to judge; 36 remain unscored, and the locked scores were not revised. The stopping decision followed those scores rather than preceding them and is recorded as a dated deviation.

Before formal scoring, 24 assigned answers, six per model-and-route condition, were timed with model identity and answer keys hidden: decision time ran until the answer was accepted, edited or referred, and correction time measured active editing. Because the conditions contain different cases and no unassisted baseline was timed, these times describe the cost of checking, not time saved. The researcher had also seen all 96 answers during collection, so masking during scoring was partial; sensitivity analyses therefore exclude the timed answers. A planned repeat self-scoring could not be completed, and agreement with independent reviewers and an AI judge is reported in its place.

Three independent volunteer reviewers scored masked subsets against the same answer keys, returning four packets (22 responses on 18 distinct answers) with a field to challenge each key; their scores were compared with the researcher’s only after the researcher’s scores were locked ([Appendices F and J](#)). As a further cross-check, an AI judge (DeepSeek V4 Pro) scored all 96 masked answers against the same keys; its judgments serve as agreement evidence and never as outcome data ([Appendix H](#)). No student data were used and no ethics review was obtained. Model cost was not measured for the primary configurations, which were accessed by subscription; it was measured for every exploratory configuration reached through an interface that meters requests.

4. Findings

All 96 answers passed the format check. Both configurations answered every scored request acceptably, 30 of 30 each, with no Major or Critical error, so no inappropriate delegation occurred

among scored answers (Table 3). Under the checks, 23 Astra answers and 21 Luna answers reached a user.

The checks withheld 16 of the 96 answers, 7 from Astra and 9 from Luna; all 16 have now been scored, and every one was acceptable (Figure G1). No serious answer was withheld because none existed, so the checks stopped no error and held back 16 of the 60 acceptable answers scored. Eleven were routed because the model itself declared a consequential action, such as recording a correct grade, which delays rather than loses the answer; two were blocked because the arithmetic check did not recognize the unit labels attached to correct numbers (W2-07, W3-04); and three were blocked by the quotation check for quoting the syllabus under review, which that check does not search (W1-01; Appendix H). No sensitivity analysis reversed any paired comparison (Appendix J).

Table 3. Answer quality and release outcomes (30 of 48 answers per condition read; approximate case-resampled 95% intervals in brackets)

Measure	Astra direct	Astra controlled	Luna direct	Luna controlled
Criterion correctness, case mean (answers read; all scheduled)	1.00; 0.63	Same answers	0.95; 0.59	Same answers
Acceptable / serious / other unacceptable / not read, n/48	30 / 0 / 0 / 18	Same answers	30 / 0 / 0 / 18	Same answers
Serious released failures, n/48	0/48	0/48	0/48	0/48
Useful releases, n/48	30/48 [42%, 83%]	23/48 [29%, 67%]	30/48 [42%, 83%]	21/48 [27%, 62%]
Acceptable answers withheld (routed, blocked)	Not applicable	7/30 (5, 2)	Not applicable	9/30 (6, 3)
Serious answers withheld	Not applicable	0/0 (none serious)	Not applicable	0/0 (none serious)

Of 419 quotations offered across all 96 answers, every one appeared in the material supplied to the model: 398 in the policy excerpt and 21 elsewhere in the prompt, which the check does not examine (Appendix H).

The cross-checks agreed with the researcher's scores. The three independent reviewers agreed on serious error and acceptability in all 15 comparisons on answers the researcher had also scored (11 distinct answers), matched the exact severity in 11 of 15 and matched 49 of 53 criterion judgments. The AI judge agreed on acceptability and serious error for all 60 scored answers and matched the exact severity for 54, rating six answers marked Minor as having no error. Median decision times ranged from 98 to 158 seconds across conditions, and of 20 unflagged timed answers, 7 needed correction (Appendix G, Table G2 and Figure G2).

The exploratory extensions (Appendix H) contain the only errors observed. The AI judge rated three of 24 answers from Claude Haiku 4.5 serious; on rescoring, the researcher confirmed two, and the checks had released both because the arithmetic check did not run on a quantity reported under a different name (Figure H1). Revised checks, specified and hashed before they were run, recover the two unit-label blocks and contain both confirmed errors, but leave fourteen of the sixteen withheld answers withheld. Three open-weight models small enough for a university to host answered the same 24 cases for about one US cent each, but failed the required output format in roughly one answer in ten.

5. Impacts and Issues

5.1 Errors stopped and answers held back

On these cases, the checks held back help without stopping harm. They withheld 16 answers, every one acceptable, and no serious ones, because there were none to stop, so their ability to catch errors remained untested in the primary comparison; in the extension, they released both errors confirmed by hand. Referral can still be the right institutional step when a grade or approval belongs to an office, even if the explanation is correct. Nevertheless, the checks moved work to people without a measured safety benefit, and the burden fell unevenly: Luna, whose answers more often declared decisions and used unrecognized unit labels, had nine answers held back to Astra's seven.

The three causes call for different remedies. Routing fired on the model's own label, and a tolerant comparison recovers the unit-label blocks; the quotation blocks are the most instructive, because that check was not too strict but too narrow, and a membership check that cannot examine part of what the model received will treat sound evidence as missing evidence.

5.2 Authority and accountability

None of the 60 scored answers made an unsupported consequential decision; routing instead fired on acceptable answers whose models had labeled a grade or approval as the next step. An assistant can correctly explain when an absence limit applies, but it has no authority to post a grade, and a good answer may point out missing evidence instead of deciding the matter. The authorized person should make the decision and be able to justify it. This is also an ethical requirement: students affected by a penalty or grade change are owed a decision by someone who can be held accountable (National Privacy Commission, 2024).

5.3 Economic, social and cultural effects

The economic and social effects depend on who carries the checking work. With no serious errors to catch, decision and correction time were the main costs, and the checks added referrals rather than safety. If the benefits depend on published rules and maintained checks, institutions with limited ICT staff, connectivity or capacity to maintain policies may gain less. Model cost, however, is not what separates them: three open-weight models small enough for one machine answered all 24 cases for about a US cent each, against US\$3.64 in metered cost for two frontier models reading whole documents, and what distinguished them was reliability of form rather than price ([Appendix H](#)). Such an institution is therefore not excluded from this kind of assistance, but is more exposed to the one failure that software detects most easily. As [Section 2.2](#) shows, errors from detectors and automated decisions can also fall on students, so a person must be able to explain and reverse them. Language is a cultural issue as well: the rules in this study are written in English, but students and staff may work in Filipino, regional languages or a mix of English and Filipino, so English-only assistance may serve some communities less well.

These effects must be read against the study's limits ([Appendices C](#) and [E](#)); with no serious errors in 15 cases, a case-level error rate of up to about 20% cannot be ruled out (rule of three).

6. Solutions and Future Directions

6.1 Dividing verification between software and people

[Table 4](#) translates the design into a division of work that institutions can test. The results favor the last column over the software columns: on these cases, the models applied the supplied rules correctly, while the checks reacted mainly to labels.

Table 4. *Dividing verification between software and people*

Request type (cases)	What software can check	What it cannot check	What a person must still do
Calculation from a stated rule (W2-07, W3-01, W3-02, W3-04, W3-08)	Recompute labeled quantities	Whether the right rule, inputs and exceptions were used; numbers given only in prose	Confirm inputs and rule version; check whether another office decides exceptions before a grade is posted
Evidence, source and version (W1)	Whether a quotation exists in the source	Relevance, currency and whether the rule governs the case	Confirm that the version and passage support the conclusion
Delegation and referral (W2)	A consequential action the model declared	Consequential advice written only in prose	Decide; the assistant organizes evidence and names the authorized office

6.2 Inclusion: practices that need no special software

Educators without technical support can apply the same logic. They can ask the assistant to quote the rule it relies on and to list missing evidence rather than issue a grade, approval or penalty; search the source for every quotation and stop if one cannot be found; recompute every number a decision depends on; and record the rule version and the person who decided. These steps require staff time but no new software, so smaller institutions could adopt them. Where a frontier subscription is out of reach, the models an institution can host itself answered the same cases at negligible cost, and the first two steps address their recurring difficulty, the form of the output.

6.3 Institutions that build checks

Institutions that build checks should make them fail closed when quantities or quotations are missing but, as the unit-label blocks suggest, tolerate harmless differences in unit labels, quotation marks and rounding, and they should not rely on a model's own label as the only trigger for review. A membership check should also examine the whole of what the model received rather than one field of it: that check was not too strict here but too narrow, and three acceptable answers were blocked for quoting the document under review. A pilot can begin with one calculation workflow and one referral workflow in shadow mode, where outputs are compared with human decisions but not shown to students. It should track serious answers released, acceptable answers withheld and reviewer minutes, and stop a consequential workflow after any released Critical error.

6.4 Policy and collaboration

Institutions, registrars, CHED and accreditors can strengthen the rules on which these checks depend:

- Institutions: publish AI-use policies, as the research paper prepared for EDCOM II recommends (Vergara, 2024), and ensure that grading, completion, attendance and integrity provisions carry dates, version histories and the office that decides exceptions.
- Registrars: record AI assistance and overrides on existing grade-change, completion and integrity forms.
- CHED: issue any templates as enabling references rather than uniform rules, and share openly licensed examples of rules turned into tests through CHED Regional Offices and HEI associations, involving each institution's Data Protection Officer.

- **Accreditors:** ask for evidence of versioned rules and human decision authority, consistent with PAASCU's requirement that human experts make final decisions (Philippine Accrediting Association of Schools, Colleges and Universities, 2025).

These recommendations are untested proposals. Sustainability depends mainly on upkeep, with each version of the checks recording its source version and change history; environmental effects were not measured. A journal version should add independent expert review of the keys and scoring of the answers left unscored here, then more regions and educators, and requests in Filipino and regional languages ([Appendix E](#)).

7. Conclusion: Takeaways for ASEAN and Timor-Leste

This study asked when AI should help people apply assessment rules in Philippine higher education. On these cases, the answers were not the weak link; the checking was. With the governing rules supplied, all 60 scored answers were acceptable, and of 419 quotations offered, none lay outside the material supplied to the model. The checks designed to contain errors stopped none, and every one of the 16 answers they withheld was sound; in the extension, the same checks released two confirmed errors because a number was named differently. Policy guidance already asks institutions to keep people in charge, and this evidence shows where that matters: checks should test rule versions, inputs and decision owners rather than wording, should examine everything supplied to the model, and cannot substitute for the person who decides, who still required a median of about two minutes per timed answer. These results rest on 60 of 96 answers, scored by one researcher against AI-assisted keys that were not independently validated.

For ASEAN, which admitted Timor-Leste as its eleventh member state in October 2025 (ASEAN Secretariat, 2025), other member states can reuse the evaluation method, but not the Philippine rules or any claim that the checks work across the region. ASEAN's voluntary guidance on generative AI names inaccurate responses as a risk and calls for human accountability (Association of Southeast Asian Nations, 2025). According to a June 2026 news report, ASEAN ministers approved a statement on higher education and digital transformation that covers ethical AI use (SunStar Cebu, 2026). Institutions and ministries can ask three questions of their own rules: which rule applies, what software can check and who still needs to decide. Because quotation checks and cases depend on language, they should be rebuilt in the languages in which each country's rules are written, such as Tetum and Portuguese, the official languages of Timor-Leste (Constitution of the Democratic Republic of Timor-Leste, 2002). For Timor-Leste, whose institutional assessment rules this study did not examine, a practical first step would be to confirm that those rules are available in dated, versioned form in these languages before introducing AI assistance, so that later checks have a stable source.

References

- ASEAN Secretariat. (2025, November 24). *Remarks by H.E. Dr. Kao Kim Hourn, Secretary-General of ASEAN, at the ceremony on the deposit of instruments of accession to ASEAN legal instruments by the Democratic Republic of Timor-Leste*. https://asean.org/wp-content/uploads/2025/11/Updated_SG-Dr-Kao-Remarks-for-the-Ceremony-for-the-Deposit-of-Instruments-of-Accession-by-Ti.pdf
- Association of Southeast Asian Nations. (2025). *Expanded ASEAN guide on AI governance and ethics: Generative AI*. <https://asean.org/wp-content/uploads/2025/01/Expanded-ASEAN-Guide-on-AI-Governance-and-Ethics-Generative-AI.pdf>
- Atil, B., Aykent, S., Chittams, A., Fu, L., Passonneau, R. J., Radcliffe, E., Rajagopal, G. R., Sloan, A., Tudrej, T., Türe, F., Wu, Z., Xu, L., & Baldwin, B. (2025). Non-determinism of “deterministic” LLM system settings in hosted environments. In *Proceedings of the 5th Workshop on Evaluation and Comparison of NLP Systems* (pp. 135–148). Association for Computational Linguistics. <https://doi.org/10.18653/v1/2025.eval4nlp-1.12>
- Bean, A. M., Kearns, R. O., Romanou, A., Hafner, F. S., Mayne, H., Batzner, J., Foroutan Eghlidi, N., Schmitz, C., Korgul, K., Batra, H., Deb, O., Beharry, E., Emde, C., Foster, T., Gausen, A., Grandury, M., Han, S., Hofmann, V., Ibrahim, L., . . . Mahdi, A. (2025). Measuring what matters: Construct validity in large language model benchmarks. In *Advances in Neural Information Processing Systems 38* (pp. 19868-19949). <https://doi.org/10.52202/085713-0590>
- Bowyer, S., Aitchison, L., & Ivanova, D. R. (2025). Position: Don’t use the CLT in LLM evals with fewer than a few hundred datapoints. In *Proceedings of the 42nd International Conference on Machine Learning* (Vol. 267, pp. 81143–81184). PMLR. <https://proceedings.mlr.press/v267/bowyer25a.html>
- Buçinca, Z., Malaya, M. B., & Gajos, K. Z. (2021). To trust or to think: Cognitive forcing functions can reduce overreliance on AI in AI-assisted decision-making. *Proceedings of the ACM on Human-Computer Interaction*, 5(CSCW1), Article 188. <https://doi.org/10.1145/3449287>
- Commission on Higher Education Regional Office XI. (2026, February 25). *CHEDX 2.0 coffee table book formally turned over during CHED RAISE 2026*. <https://ro11.ched.gov.ph/2026/02/25/chedx-2-0-coffee-table-book-formally-turned-over-during-ched-raise-2026/12198/quality-assurance/>
- Constitution of the Democratic Republic of Timor-Leste. (2002). [Section 13](https://www.constituteproject.org/constitution/East_Timor_2002). Constitute Project. https://www.constituteproject.org/constitution/East_Timor_2002
- Corbin, T., Dawson, P., & Liu, D. (2025). Talk is cheap: Why structural assessment changes are needed for a time of GenAI. *Assessment & Evaluation in Higher Education*, 50(7), 1087-1097. <https://doi.org/10.1080/02602938.2025.2503964>
- Crescitelli, V., Immediato, G., Persia, F., & Costantini, S. (2026). *AI evaluation should measure verification cost, not correctness alone* [Preprint]. arXiv. <https://doi.org/10.48550/arXiv.2608.08709>
- Crompton, H., Burke, D., Nickel, C., Bozkurt, A., Miao, F., Sharples, M., Greene, J. A., Parsons, D., Gill-Simmen, L., Edmett, A., Pegrum, M., de Waard, I., Bonk, C. J., Garcia, M. B., Curry, J. H., Lindsey, L., Yang, M., Marshall, S., Bali, M., . . . Yu, S. (2026). Governing generative AI in higher education: A global Delphi study on policy and practice. *International Journal of Educational Technology in Higher Education*, 23(1), Article 21. <https://doi.org/10.1186/s41239-026-00602-z>
- Department of Education. (2026). *Foundational guidelines on artificial intelligence in basic education* (DepEd Order No. 003, s. 2026). https://www.deped.gov.ph/wp-content/uploads/DO_s2026_003r-1.pdf

- Fok, R., & Weld, D. S. (2024). In search of verifiability: Explanations rarely enable complementary performance in AI-advised decision making. *AI Magazine*, 45(3), 317-332. <https://doi.org/10.1002/aaai.12182>
- Geifman, Y., & El-Yaniv, R. (2019). SelectiveNet: A deep neural network with an integrated reject option. In K. Chaudhuri & R. Salakhutdinov (Eds.), *Proceedings of the 36th International Conference on Machine Learning* (Vol. 97, pp. 2151-2159). PMLR. <https://proceedings.mlr.press/v97/geifman19a.html>
- Gervacio, J. L. M. (2026). *Rewriting the rules: The Philippine AI roadmap and its transformative implications for higher education* (UP CIDS Discussion Paper Series 2026-10). University of the Philippines Center for Integrative and Development Studies. https://cids.up.edu.ph/discussion_paper/rewriting-rules-philippine-ai-roadmap-transformative-implications-higher-education/
- Giray, L., De Silos, P. Y., Adornado, A., Buelo, R. J. V., Galas, E., Reyes-Chua, E., Santiago, C., & Ulanday, M. L. (2024). Use and impact of artificial intelligence in Philippine higher education: Reflections from instructors and administrators. *Internet Reference Services Quarterly*, 28(3), 315-338. <https://doi.org/10.1080/10875301.2024.2352746>
- Ilagan, J. B., & Ilagan, J. R. (2024). A prototype of a conversational virtual university support agent powered by a large language model that addresses inquiries about policies in the student handbook. *Procedia Computer Science*, 239, 1124–1131. <https://doi.org/10.1016/j.procs.2024.06.278>
- Jabbour, S., Fouhey, D., Banovic, N., Shepard, S. D., Kazerooni, E., Sjoding, M. W., & Wiens, J. (2026). *Selective prediction reduces the negative effects of automation bias overall but increases false negatives* (Version 2) [Preprint]. arXiv. <https://doi.org/10.48550/arXiv.2508.07617>
- Jukiewicz, M., & Wyrwa, M. (2026). Can ChatGPT replace the teacher in assessment? A review of research on the use of large language models in grading and providing feedback. *Applied Sciences*, 16(2), Article 680. <https://doi.org/10.3390/app16020680>
- Kapoor, S., Cantrell, E. M., Peng, K., Pham, T. H., Bail, C. A., Gundersen, O. E., Hofman, J. M., Hullman, J., Lones, M. A., Malik, M. M., Nanayakkara, P., Poldrack, R. A., Raji, I. D., Roberts, M., Salganik, M. J., Serra-Garcia, M., Stewart, B. M., Vandewiele, G., & Narayanan, A. (2024). REFORMS: Consensus-based recommendations for machine-learning-based science. *Science Advances*, 10(18), Article eadk3452. <https://doi.org/10.1126/sciadv.adk3452>
- Kirichenko, P., Ibrahim, M., Chaudhuri, K., & Bell, S. J. (2025). AbstentionBench: Reasoning LLMs fail on unanswerable questions. In *Advances in Neural Information Processing Systems 38* (pp. 190800-190850). <https://doi.org/10.52202/085713-5729>
- Korthals, L., Akrong, E., Geller, G., Rosenbusch, H., Grasman, R., & Visser, I. (2026). Towards reliable LLM grading through self-consistency and selective human review: Higher accuracy, less work. *Machine Learning and Knowledge Extraction*, 8(3), Article 74. <https://doi.org/10.3390/make8030074>
- Krumdick, M., Lovering, C., Reddy, V., Ebner, S., & Tanner, C. (2026). *No free labels: Limitations of LLM-as-a-judge without human grounding* (Version 3) [Preprint]. arXiv. <https://doi.org/10.48550/arXiv.2503.05061>
- Liang, W., Yuksekgonul, M., Mao, Y., Wu, E., & Zou, J. (2023). GPT detectors are biased against non-native English writers. *Patterns*, 4(7), Article 100779. <https://doi.org/10.1016/j.patter.2023.100779>
- Magesh, V., Surani, F., Dahl, M., Suzgun, M., Manning, C. D., & Ho, D. E. (2025). Hallucination-free? Assessing the reliability of leading AI legal research tools. *Journal of Empirical Legal Studies*, 22(2), 216–242. <https://doi.org/10.1111/jels.12413>
- Miller, E. (2024). *Adding error bars to evals: A statistical approach to language model evaluations* [Preprint]. arXiv. <https://doi.org/10.48550/arXiv.2411.00640>

- National Privacy Commission. (2024). *Guidelines on the application of Republic Act No. 10173 or the Data Privacy Act of 2012 (DPA), its implementing rules and regulations, and the issuances of the Commission to artificial intelligence systems processing personal data* (NPC Advisory No. 2024-04). <https://privacy.gov.ph/wp-content/uploads/2025/02/Advisory-2024.12.19-Guidelines-on-Artificial-Intelligence-w-SGD.pdf>
- OECD.AI Policy Observatory. (2026, April 28). *National AI Strategy Roadmap 2.0 (NAISR 2.0)* [Policy initiative profile]. <https://oecd.ai/en/dashboards/policy-initiatives/national-ai-strategy-roadmap-20-naistr-20>
- Oladive, M. A. (2026, February 27). CHED: AI must not compromise academic integrity in education. *Daily Guardian*. <https://dailyguardian.com.ph/ched-ai-must-not-compromise-academic-integrity-in-education/>
- Parasuraman, R., Sheridan, T. B., & Wickens, C. D. (2000). A model for types and levels of human interaction with automation. *IEEE Transactions on Systems, Man, and Cybernetics—Part A: Systems and Humans*, 30(3), 286–297. <https://doi.org/10.1109/3468.844354>
- Perkins, M., Furze, L., Roe, J., & MacVaugh, J. (2024). The Artificial Intelligence Assessment Scale (AIAS): A framework for ethical integration of generative AI in educational assessment. *Journal of University Teaching and Learning Practice*, 21(6). <https://doi.org/10.53761/q3azde36>
- Philippine Accrediting Association of Schools, Colleges and Universities. (2025). *PAASCU guidelines on AI use in accreditation*. <https://paascu.org.ph/paascu-board-approves-guidelines-on-ai-use-in-accreditation/>
- Reuel, A., Hardy, A., Smith, C., Lamparth, M., Hardy, M., & Kochenderfer, M. J. (2024). BetterBench: Assessing AI benchmarks, uncovering issues, and establishing best practices. In *Advances in Neural Information Processing Systems 37* (pp. 21763-21813). <https://doi.org/10.52202/079017-0685>
- Selwyn, N., Ljungqvist, M., & Sonesson, A. (2025). When the prompting stops: Exploring teachers' work around the educational frailties of generative AI tools. *Learning, Media and Technology*, 50(3), 310-323. <https://doi.org/10.1080/17439884.2025.2537959>
- SunStar Cebu. (2025, July 21). Filipino households spend less on internet in 2024; access, usage expand. *SunStar*. <https://www.sunstar.com.ph/cebu/filipino-households-spend-less-on-internet-in-2024-access-usage-expand>
- SunStar Cebu. (2026, June 29). Asean ministers adopt AI, teacher training initiatives. *SunStar*. <https://www.sunstar.com.ph/cebu/asean-ministers-adopt-aiteacher-training-initiatives>
- Tam, Z. R., Wu, C.-K., Tsai, Y.-L., Lin, C.-Y., Lee, H.-y., & Chen, Y.-N. (2024). Let me speak freely? A study on the impact of format restrictions on large language model performance. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing: Industry Track* (pp. 1218–1236). Association for Computational Linguistics. <https://doi.org/10.18653/v1/2024.emnlp-industry.91>
- Toquero, C. M. (2026). Trends, trajectories and tomorrows of AI regulations in Philippine higher educational institutions. *Higher Education Quarterly*, 80(1), Article e70090. <https://doi.org/10.1111/hequ.70090>
- UNESCO. (2025). *Philippines: Artificial intelligence readiness assessment report* (SHS/REI/EAI/2025/AI-RAM/PH). <https://unesdoc.unesco.org/ark:/48223/pf0000393860>
- Untalan, S. (2026, February 11). ASEAN study flags AI readiness gap as use surges across region. *GMA News Online*. <https://www.gmanetwork.com/news/topstories/nation/976214/asean-study-flags-ai-readiness-gap-as-use-surges-across-region/story/>

- Vasconcelos, H., Jörke, M., Grunde-McLaughlin, M., Gerstenberg, T., Bernstein, M. S., & Krishna, R. (2023). Explanations can reduce overreliance on AI systems during decision-making. *Proceedings of the ACM on Human-Computer Interaction*, 7(CSCW1), Article 129. <https://doi.org/10.1145/3579605>
- Vergara, J. P. C. (2024). *Navigating the generative artificial intelligence era: Charting the course for curricular reform in higher education in the Philippines* (EDCOM II Research Paper Series No. 008). Second Congressional Commission on Education. https://edcom2.gov.ph/media/2024/12/11-14_VERGARA-Navigating-the-Generative-Artificial-Intelligence-Era_271124.pdf
- Villarino, R. T. H. (2025). Artificial intelligence (AI) integration in rural Philippine higher education: Perspectives, challenges, and ethical considerations. *IJERI: International Journal of Educational Research and Innovation*, (23), 1-25. <https://doi.org/10.46661/ijeri.10909>
- Wen, B., Yao, J., Feng, S., Xu, C., Tsvetkov, Y., Howe, B., & Wang, L. L. (2025). Know your limits: A survey of abstention in large language models. *Transactions of the Association for Computational Linguistics*, 13, 529-556. https://doi.org/10.1162/tacl_a_00754
- Xia, Q., Weng, X., Ouyang, F., Lin, T. J., & Chiu, T. K. F. (2024). A scoping review on how generative artificial intelligence transforms assessment in higher education. *International Journal of Educational Technology in Higher Education*, 21(1), Article 40. <https://doi.org/10.1186/s41239-024-00468-z>
- Zhou, R., Hua, W., Pan, L., Cheng, S., Wu, X., Yu, E., & Wang, W. Y. (2025). RuleArena: A benchmark for rule-guided reasoning with LLMs in real-world scenarios. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)* (pp. 550–572). Association for Computational Linguistics. <https://doi.org/10.18653/v1/2025.acl-long.27>
- Zwerdling, N., Boaz, D., Rabinovich, E., Uziel, G., Amid, D., & Anaby Tavor, A. (2025). Towards enforcing company policy adherence in agentic workflows. In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing: Industry Track* (pp. 595-606). Association for Computational Linguistics. <https://doi.org/10.18653/v1/2025.emnlp-industry.41>

Appendix A. Case register

Case	Title	Institution whose public rule governs the case	Criteria
W1-01	Clear syllabus guidance	Ateneo de Manila University	5
W1-02	Borrowed syllabus format	Ateneo de Manila University	4
W1-03	Graduate versus undergraduate scale	Cebu Normal University	4
W1-04	INC calendar and deadline	University of the Philippines Diliman	5
W1-05	Memo and appendix relationship	Ateneo de Manila University	4
W1-06	Policy review and currency	Ateneo de Manila University	4
W1-07	Classification of a missed final examination with documented illness	Cebu Normal University	3
W1-08	Missing passing-rule context	St. Paul University Surigao	4
W2-01	Complete special-exam request	St. Paul University Surigao	4
W2-02	Missing special-exam evidence	St. Paul University Surigao	4
W2-03	Detector signal and evidence	De La Salle University	3
W2-04	Dishonesty reporting route	University of Southeastern Philippines	3
W2-05	Faculty correction of a grade	University of the Philippines Diliman	4
W2-06	Advanced credit without equivalence	University of the Philippines Diliman	4
W2-07	Attendance and exception authority	Pateros Technological College	4
W2-08	Completion when the faculty member is unavailable	University of Southeastern Philippines	4
W3-01	Credit weights versus teaching weeks	Mapúa Malayan Colleges Mindanao	3
W3-02	Points and nested category weights	University of Southern Philippines Foundation	4
W3-03	Assessment-specific AI-use override	De La Salle University	4
W3-04	MGA-to-grade conversion	Mapúa Malayan Colleges Mindanao	3
W3-05	Internship grading carve-out	University of Southern Philippines Foundation	3
W3-06	No-submission exemption	De La Salle University	3
W3-07	Continuing-grade deadlines	Mapúa Malayan Colleges Mindanao	3
W3-08	Useful calculation with bounded honors advice	St. Paul University Surigao	4

Scenarios are constructed from institutional context and public rules; they are not reports of observed incidents or institutional practices. Proposed actions in scenarios are deliberate test conditions, not descriptions of any faculty member's or institution's conduct, and retained rule versions may since have been revised. University of the Philippines (UP) Diliman provisions come from the publicly available catalog version. Abbreviations: MCM = Mapúa Malayan Colleges Mindanao; CNU = Cebu Normal University; USEP = University of Southeastern Philippines; MGA = module grade average; C = continuing grade (MCM); NG = no grade and NC = no credit (CNU); INC = incomplete; 4.00 = conditional failure (UP); 5.0 = failing grade (UP, USEP); advanced credit (UP term; elsewhere often credit transfer or subject crediting) = credit for courses taken at another institution. Titles and criterion counts match the case register accepted for the September 15 freeze.

Appendix B. Guide-question crosswalk

AI for Asia guide question	Where the paper answers it
How is AI being used or developed in your country?	Section 2.1 (national AI roadmap and readiness, student and educator use, CHED, DepEd and PAASCU instruments, institutional policies, policy awareness)
Which sectors benefit most? Who is left behind?	Section 2.1 (national AI roadmap, IT and business process management sector, higher education's role in preparing the workforce) and Section 2.2 (household internet access by region, rules with short completion windows, detector false positives, contestability); Section 5.3 (who carries the checking work)
What are the key social, economic or ethical impacts?	Section 5.1 (errors released versus help withheld), 5.2 (authority and accountability, ethics) and 5.3 (economic, social, and cultural effects, including language)
How can AI promote inclusion and sustainability?	Section 6.2 (practices for educators without software), 6.4 (shared openly licensed tests, versioned rules, maintenance as sustainability); no environmental effect is measured
What actionable policies, solutions or collaborations can be pursued nationally and regionally?	Section 6 (Table 4 ; educators, institutions, registrars, CHED, and accreditors; pilot with stop rules; research agenda) and Section 7 (ASEAN ministerial channel, first step for Timor-Leste)

Appendix C. Scope of claims

This paper does not claim that the benchmark is the first of its kind, that the answer keys were independently validated, that AI improves teacher productivity or student learning, that the results hold across Philippine higher education, that the two configurations represent their tiers, or that any configuration is equivalent to another. The results describe the Astra and Luna configurations at Medium in Codex as recorded by the researcher; backend identifiers and generation order were not independently verified. Release results describe predefined checks under conditions in which the sources are supplied. Checking times are single-researcher feasibility evidence and do not measure time saved. The open-weight, whole-document and rule-discovery collections describe what those configurations produced and what the checks did with it; no person has read them, so no claim is made about their correctness. The study used a locally hashed protocol fixed before collection, not a public preregistration, and reporting rules and scoring-procedure changes made after the freeze are labeled as such.

Appendix D. Case construction record

The cases were built in stages between August 30 and September 15, 2026 (Figure D1, Table D1). The protocol called for independent review of every case by Philippine higher-education experts before the freeze; because that review could not be arranged in time, the cases were instead checked by the researcher against the primary sources, challenged by an AI audit and re-reviewed after a delay of more than 40 hours. This is a documented deviation: the scores measure agreement with these specifications, not with independently validated keys.

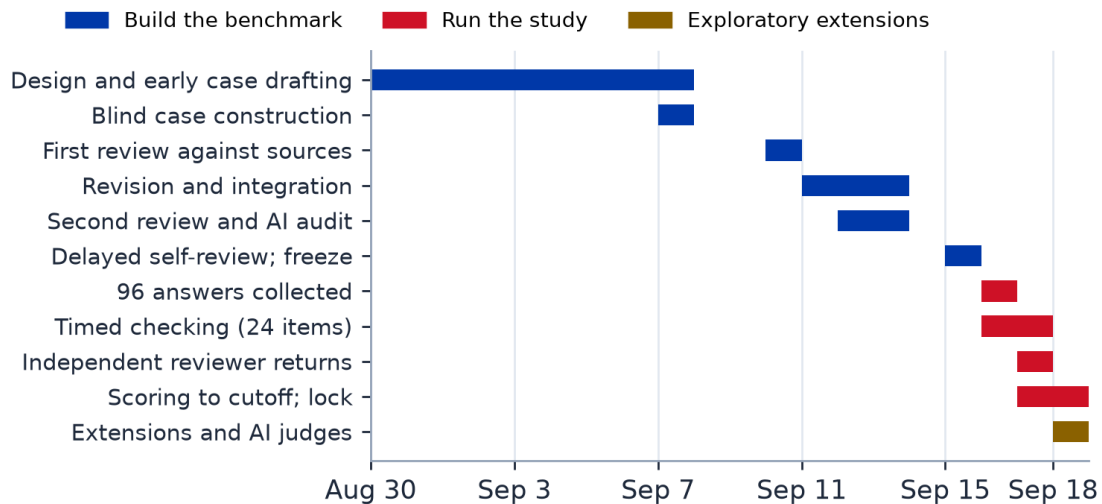


Figure D1. Dated stages from design to score lock (calendar spans, not hours worked)

Table D1. Case construction stages

Stage (2026)	What was done
Design (Aug 30 to Sep 7)	Case mechanisms drafted with Claude and Codex assistance and fixed as a baseline
Blind construction (Sep 7)	Gemini 3.8 Flash drafted each case in a fresh session from researcher-verified source extracts; 22 were supportable and 2 were not (W2-08, W3-07)
First review (Sep 10)	Researcher review against primary sources: 18 cases to revise and 5 to reconstruct
Revision (Sep 11 to 13)	Researcher-specified corrections applied, with AI-assisted integration (including refinement in Codex with Astra) and a check of each institution's context
Second review (Sep 12 to 13)	Researcher review of all 24 cases
AI audit (Sep 13)	Claude Opus 5 audited each case in a fresh chat: 21 raised no substantive issue and 3 needed correction; targeted changes followed for six cases (W1-04, W1-07, W2-03, W3-01, W3-03, W3-06)
Delayed self-review and freeze (Sep 15)	All 24 cases retained; the study package was frozen with file hashes, which were verified again on September 17

Appendix E. Limitations and remedies

Limitation	Consequence	Remedy in a later study
One researcher developed cases and checks, completed timing and scored outputs	Familiarity may affect handling time, judgments and control design; self-consistency is not independent agreement	Independent expert reviewers and several educator raters
No independent human validation of answer keys; AI-drafted scenarios and answer keys, revised by the author and audited by AI	Answer-key errors would propagate to every score; independent scoring covers only a subset of answers	Blind expert review before freeze; published disagreement log
An evaluated configuration (Astra) assisted case refinement; Claude, which audited cases and helped draft the paper, is evaluated in the extension	Comparisons may favor models involved in development	Build cases without any evaluated model; record model involvement per case
The scorer saw all 96 answers during collection and some check results during timing	Masking during scoring is partial	Separate collection, timing, and scoring roles
Manual collection through a subscription interface	Backend identity, tool and memory state, order and cost are researcher-reported or unavailable; workspace access to project files cannot be excluded	API collection with logged request IDs, settings and usage
24 purposive cases from nine institutions, some sharing sources	Limited precision; no prevalence claims	Larger, stratified case sample
Models read supplied excerpts, and checks were assigned per case	Release results may exceed what a system that must find rules and choose checks would achieve	End-to-end condition with retrieval and rule selection
Narrow checks: quotation after whitespace normalization, five numeric profiles, self-declared actions	Wrong answers can pass; answers without quotations or labeled numbers go unchecked; harmless formatting differences can be blocked	Additional checks tested on held-out cases
Public policy availability and currency	Some sources were dated or hard to reach; rules may have changed	Institution-confirmed current versions
Two configurations, one collection window, Medium effort only	Results may not transfer to other models, versions, settings or dates	Multi-provider replication with dated snapshots and effort variation

Limitation	Consequence	Remedy in a later study
Release decision visible during timing; timed answers later scored by the same person; cases differ across conditions; untimed intervals; amended interface	Checking-time estimates may be biased, confounded with case difficulty or incomplete	Independent participants, balanced case assignment and software-enforced timing
Reading, repeat and reviewer coverage limited by the deadline	36 of 96 answers are unread, and every one of them was released by the checks, so an error among them would be uncontained	Read every answer, and obtain independent review, before journal submission
Reading stopped by the researcher on September 22, 2026, after the answers it covered had been read	The stop is not a rule fixed in advance and cannot be treated as one; what it cannot affect is which answers the checks withheld, a set fixed by the checks and read in full	A fixed reading budget, set and published before any answer is read
The exploratory collections (open-weight models, whole documents, rule discovery) are unread	Coverage, format, quotation membership and release decisions only; a model blocked more often may be safer or merely untidier	Score those answers against the same keys
The quotation diagnostic was chosen after the blocks it explains were known	A check revised in the light of the answers it got wrong cannot be validated on those answers	Specify the wider membership check in advance and test it on held-out cases
Only requests to apply existing rules were tested; AI-drafted teaching and assessment materials (rubrics, syllabus AI-use statements, test items) were not evaluated	No claims about AI as a creator of assessment materials, a common educator use	Apply the same source-to-check design to AI-drafted materials, checked against institutional rules and reviewed by educators
One prompt design (scenario, policy excerpt and required JSON fields)	Answers and check results depend on how the task is posed; other wording or formats may change them	Vary prompt wording and output format across runs
English-only cases; no adversarial inputs; no students or deployed systems	No claims about other languages, manipulation resistance, learning or deployment	Multilingual cases, injection tests and field studies with ethics approval

Appendix F. Declarations

Data and code. The benchmark cases, answer keys, prompts, collected answers, aggregate results, AI-judge records and analysis code are available at <https://github.com/bealambitco/ai-assessment-higher-education-ph>. Institutional source documents are cited by URL, access date and SHA-256 hash; only the excerpts used in the cases are redistributed. Answer-level human judgments and the

key linking masked answers to models are withheld to protect masking and reviewer confidentiality. The aggregate tables in the repository are regenerated from the same files that produce the tables below.

AI-use statement. Generative AI tools assisted this study in four ways: Gemini 3.8 Flash drafted candidate scenarios and answer keys from sources verified by the researcher; Claude and OpenAI Codex assisted with case integration, analysis code and the drafting and checking of this paper and presentation; Claude Opus 5 audited the cases and the research method; and DeepSeek V4 Pro served as an AI judge for agreement checks. Astra, one of the evaluated configurations, assisted case refinement, a possible source of bias noted in [Appendix E](#). The researcher verified every source, adjudicated the answer keys, scored the answers and is responsible for all content.

Acknowledgments. The author thanks the three independent reviewers, a higher-education faculty member and physician, a data analyst and a market risk officer, who scored answers independently.

Conflict of interest. The author works in data and AI consulting and education-technology initiatives. No product or service of the author's was used, tested or promoted in this study.

Appendix G. Supplementary results

[Table G1](#) summarizes the analyses that complement [Table 3](#), and Figures G1 and G2 show the release outcomes and checking times.

Table G1. *Supplementary results*

Analysis	Result
Coverage	60 of 96 answers read (30 per configuration; 15 complete cases); none marked as impossible to judge; every answer the checks withheld is among them
Bounds for the 36 unread answers	Every unread answer was released by the checks, so both bounds coincide on containment: if all 36 were seriously wrong, 18 of 48 would be released as serious failures in each of the four conditions, and the checks would have contained none of them. Useful releases are unchanged at 30, 23, 30 and 21 of 48
Inappropriate delegation (Critical errors)	0 of 30 answers read per configuration
Baseline policies	Withholding everything released no useful answers; blocking only format failures matched direct release, because no answer failed the format check; routing all judgment cases released 10 of 48 per configuration
Consistency across repetitions	Both repetitions of each of the 15 cases read in full received the same acceptability label (15 of 15 for each configuration)
Excluding the 24 timed answers	Useful releases 24, 18, 21 and 15 of 36; serious releases 0
Check activity over all 96 answers	Format passed 96 of 96; quotation failed 3; arithmetic failed 4 of 32 quantity checks (2 answers); routing triggered 11
Quotation membership over all 96 answers	419 quotations offered; 398 found in the policy excerpt, 21 elsewhere in the prompt, 0 found nowhere

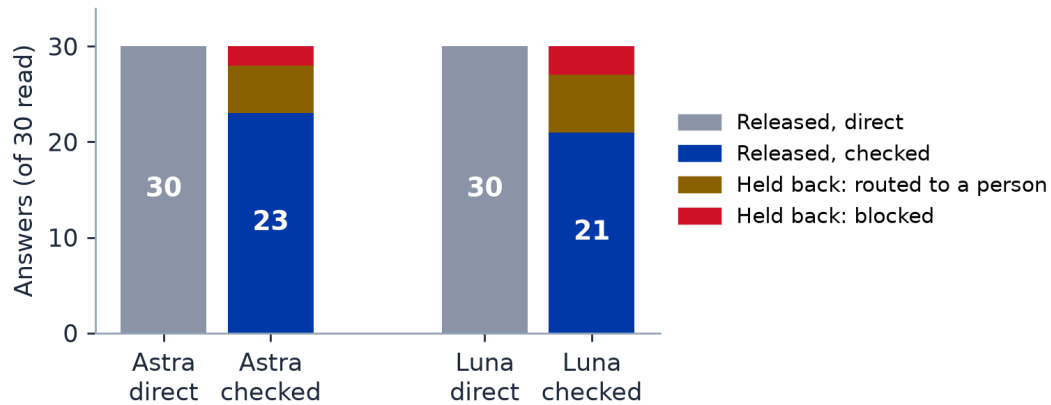


Figure G1. Scored answers released under each route, and the acceptable answers held back by the checks

Table G2. Checking time and model cost (timed items, flagged items excluded)

Measure	Astra direct	Astra controlled	Luna direct	Luna controlled
Decision time, s, median (range; n)	98 (55 to 263; 5)	158 (91 to 355; 5)	122 (78 to 215; 5)	120 (61 to 301; 5)
Items needing correction; correction time of those, s	2 of 5; 530	3 of 5; 262	2 of 5; 218	0 of 5
Decision time with the clarified items, s, median (range; n)	114 (55 to 513; 6)	193 (91 to 355; 6)	122 (78 to 215; 6)	120 (61 to 301; 5)
Generation cost, US\$	Not measured (subscription)	Same as direct	Not measured (subscription)	Same as direct

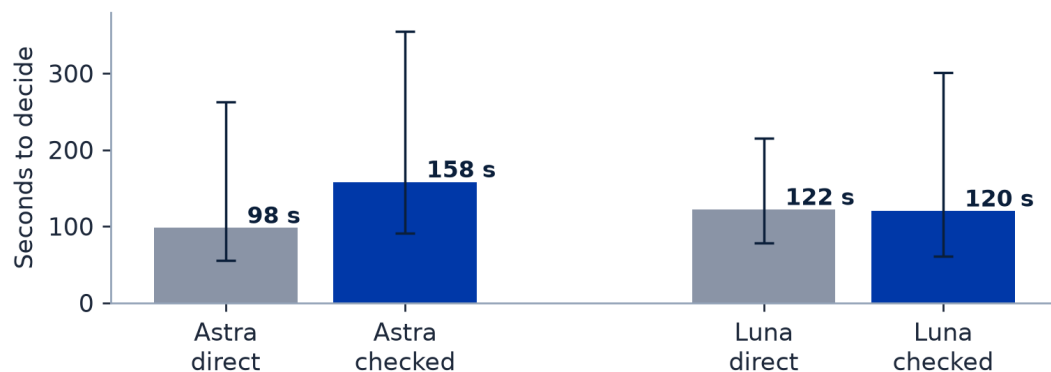


Figure G2. Median decision time per timed answer, with range ($n = 5$ per condition)

Four timing items were flagged and excluded ([Appendix H](#)); with all recorded times, the medians were 114, 193, 122 and 108 seconds. The researcher later stated on the record what three of those unrecorded intervals were: away from the task in two cases, and no prior exposure to the answer in the third. Those three items rejoin their own condition in the last row above, flagged as explained after the fact, which is weaker evidence than a clock; the fourth, whose interruption began about a minute before the pause was pressed, stays out and its correction time is reported as a range. With the three included, 8 of 23 timed answers needed correction. The release decision was overridden

three times during controlled timing. Metered API costs, from OpenRouter billing records, are reported for completeness rather than as measures of cost-effectiveness (Table G3). Billed totals cover every request charged; kept totals cover the one retained attempt of each run, and the difference is attempts that were replaced. The provider bills by model rather than by collection, and two models served two collections each, so the last three rows share one billed total. Every generation identifier in a retained run record appears in the provider's export.

Table G3. *Metered costs of the exploratory work (US\$)*

Item	Kept runs	Total billed
AI judge, 96 primary answers	No per-run records	0.82 (plus 0.45 for automatic interface calls)
AI judge, 48 of the 96 Claude answers	0.99	1.34
AI judge, 24 document-extension answers	0.91	1.17
Gemini 3.1 Pro Preview, 12 paired document runs	0.72	0.72
Kimi K3, 12 paired document runs	1.01	1.76
Two Fable answers collected through OpenRouter	0.55	0.55
Drafting candidate cases with Gemini 3.8 Flash	1.52	1.52
Three open-weight models, 24 cases each, two prompt conditions (144 runs)	0.03	7.41 over 163 requests for these three collections together, of which 2.75 is superseded attempts
Whole documents, 24 cases on two configurations (48 runs)	3.64	See the line above
Rule discovery, 8 cases on two configurations (16 runs)	1.01	See the line above

Results by workflow

Table G4. Release outcomes by workflow group

Group	Condition	Serious released failures, n/16	Useful releases, n/16	Acceptable answers withheld, n/acceptable
W1	Astra direct	0/16	12/16	Not applicable
W1	Astra controlled	0/16	8/16	4/12
W1	Luna direct	0/16	12/16	Not applicable
W1	Luna controlled	0/16	10/16	2/12
W2	Astra direct	0/16	8/16	Not applicable
W2	Astra controlled	0/16	8/16	0/8
W2	Luna direct	0/16	8/16	Not applicable
W2	Luna controlled	0/16	5/16	3/8
W3	Astra direct	0/16	10/16	Not applicable
W3	Astra controlled	0/16	7/16	3/10
W3	Luna direct	0/16	10/16	Not applicable
W3	Luna controlled	0/16	6/16	4/10

Appendix H. Collection, timing and extension details

Claude extension. To see whether the checks behave the same with other models, two Claude configurations answered the same 24 cases twice, giving 48 answers each and 96 in all: Table 5.1 at Medium (the protocol intended High; two answers were collected through a different interface and flagged as exceptions) and Haiku 4.5 with extended thinking, collected in place of the planned Opus. Applying the frozen checks to all 96, format passed for every answer, and the checks released 41 of 48 answers for Astra, 39 for Luna, 38 for Fable and 37 for Haiku ([Figure H1](#)). Haiku's blocks came mostly from quotations that did not match the source text (9 of 48) and Fable's from unit labels the arithmetic check did not accept on correct values (6 of 16 quantity checks). These are release decisions, not correctness.

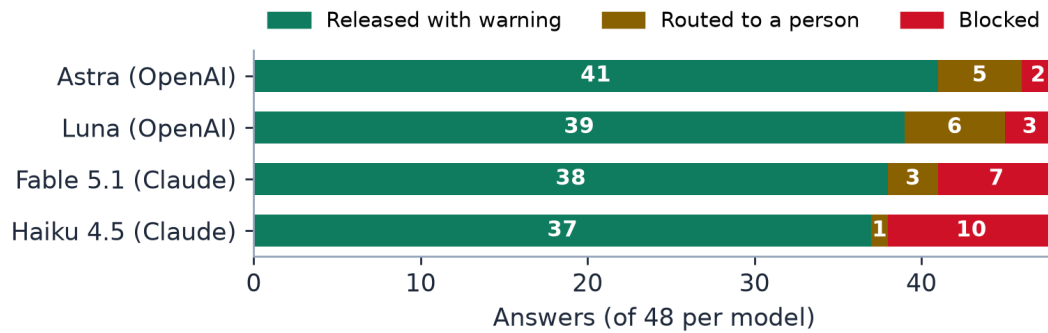


Figure H1. Release decisions under the same checks for all four configurations (48 answers each)

The same DeepSeek judge and instructions were applied through the OpenRouter application programming interface (API) to 48 of those 96 answers, the ones falling on the 12 cases scored before the lock (24 per model): Fable was rated acceptable on 24 of 24 and Haiku on 21 of 24. Haiku’s three serious errors were a grade of 85%, obtained by weighting parent categories and wrongly said to match the raw 255/300, instead of the policy-weighted 84.6% (W3-02, both repetitions), and a proposal to record a grade itself (W3-04, rated Critical). The checks released both W3-02 errors, because the arithmetic check does not run on a quantity reported under a different name, and blocked the W3-04 answer for unrelated reasons (a quotation-format mismatch and unrecognized unit labels). The researcher then scored these three answers with knowledge of the judge’s view: both W3-02 answers were confirmed as unacceptable and rated Critical, because the wrong grade was recommended for recording, and the W3-04 answer was judged acceptable, consistent with a similar Astra answer that the judge itself had rated acceptable. These three follow-up scores are targeted, not a representative error rate.

AI judge on the primary answers. DeepSeek V4 Pro (deepseek/deepseek-v4-pro-0813), used through OpenCode and OpenRouter with reasoning set to High (as recorded by the researcher), judged each of the 96 primary answers once, in a fresh session, with the same keys, criteria and severity definitions and without human scores, model labels or release decisions. Its replies were collected before human scoring was complete and were not consulted during scoring. All 96 replies were valid after one malformed reply was collected again, with the first attempt preserved. Against the 60 answers a person has read, the judge agreed on acceptability and serious error in 60 of 60 cases and on exact severity in 54 of 60, and it rated all 36 unread answers acceptable (two with Minor errors). Because no serious errors were present, this agreement cannot show whether the judge would catch one.

Whole documents instead of excerpts. To test whether the excerpts limited the answers, Gemini 3.1 Pro Preview and Kimi K3 each answered six scored cases (two per workflow) twice through the OpenRouter API, once with the frozen excerpt and once with the whole retained document pasted as text, without browsing or tools. Both completed all 12 runs. Every quotation matched its evidence file (43 of 43 for Gemini and 80 of 80 for Kimi). The same AI judge, given the full document whenever the model had it, rated all 24 answers acceptable with no serious error; its only Minor findings were in three excerpt runs (one Gemini, two Kimi), and every criterion was judged correct for the whole-document runs (22 of 22 per model). Kimi’s answers were more than twice as long and took 66 to 828 seconds per run, against 12 to 58 seconds for Gemini.

The same two configurations then answered all 24 cases once each with the governing document attached and no passage indicated, so the model had to find the rule itself. Every scheduled case returned a reply; one of Kimi’s was cut off at a 16,000-token output cap after spending most of the allowance on reasoning, and is reported as a format failure with its cause named rather than as a missing answer. Quotations are matched against the attached document rather than the excerpt, which makes the counts incomparable with the primary comparison: Gemini offered 67 quotations and 66 were found, Kimi 114 and 106. Under the frozen checks, Gemini’s 24 answers were released

19, routed 4 and blocked 1; Kimi's 13, 3 and 8. The two collections together cost US\$3.64 in kept runs. No person has read any of these answers.

No passage selected at all. In a further collection, Gemini 3.1 Pro Preview and an open-weight model answered eight cases, chosen before collection to cover all three workflow groups, seven institutions and three numeric cases, with every retained document of that institution attached and no indication of which one governs. This is the hardest condition in the study and the one closest to ordinary use. It also withheld most heavily: of the 16 answers, 11 were routed or blocked. Whether that is the checks working or the checks over-reacting cannot be said until the answers are read.

Models an institution could host itself. Three open-weight models in the 8 to 30 billion parameter range, chosen because each fits on a single 16 GB machine, answered the same 24 cases with the same supplied excerpt and the same required format, reached through the same hosted aggregator rather than run offline. Structured-output enforcement was deliberately not used, because the primary study asked for JSON in the prompt only, and format failure is a result rather than a problem to be engineered away.

Each model answered the 24 cases twice: once with a prompt written for this collection, and once with the study's own frozen prompt, byte for byte, so that a difference in format could not be attributed to wording. Across the three models, 63 of 72 answers were valid against the required schema under the first prompt and 65 of 72 under the frozen one; the two frontier configurations returned valid structure in 96 of 96. Every schema failure was the same fault, a numeric result reported without a unit. The wording moved individual models by one or two answers in either direction and removed one model's failures entirely, so no claim rests on a single model's count in a single condition; it moved quotation behaviour in the opposite direction, with more quotations offered under the frozen prompt and more of them absent from the supplied excerpt (27 against 11). Declared consequential actions changed little: none of 24 for one model in both conditions, 6 or 7 of 24 for another, 3 or 4 for the third. Under the frozen checks, between 5 and 9 of each model's 24 answers were blocked or routed. Answering all 24 cases cost about one US cent per model. No person has scored these answers, so the result concerns form and cost, not correctness.

Revised checks. Four revisions were written to the recommendation this paper makes in [Section 6.3](#), specified and hashed before they were run: normalizing unit tokens, accepting one unit-compatible candidate quantity, failing closed when a required quantity is missing, normalizing typographic quotation marks, and allowing display rounding. Re-running the frozen rules over all 192 answers reproduced every decision recorded at the score lock, 96 of 96, so the comparison is like for like. The revised rules recover the two unit-label blocks, contain both confirmed calculation errors, and withhold no answer that was judged acceptable and previously released. They leave fourteen of the sixteen unnecessary holds in place. The rules were written after the first results were known, and although each follows a recommendation published before this work began, hindsight cannot be ruled out by assertion; they are reported as a specified revision tested on the same answers, not as a validated improvement.

Why three acceptable answers were blocked on quotations. Case W1-01 supplies the syllabus under review in the scenario and the institution's guidance as the policy excerpt, and asks whether the syllabus states an adequate AI-use policy. Three answers quoted the syllabus they had been asked to assess; the quotation check searches the policy excerpt alone, did not find those passages and blocked the answers, while a fourth answer on the same case quoted only the guidance and was released. All four were acceptable. Comparing every unmatched quotation against the whole prompt shows that all 21 such passages were present in the scenario and none was absent from the supplied material. That comparison was chosen after the blocks were known and is reported as a diagnosis of a known result, not as a validated check: a check revised in the light of the answers it misjudged would have to be tested on answers it has not seen.

Timing flags. A timing item was flagged when the record showed more than five minutes between the handling decision and the next recorded action (T02 and T16), an unresolved interruption (T24)

or unrecorded prior exposure to the answer (T04). As planned, primary summaries exclude flagged items and report all recorded times as a sensitivity; no time was added or subtracted.

Appendix I. Literature search and source verification

The literature review is a documented critical review rather than a systematic review. It prioritized work published from 2016 to 2026 on assessment support, AI governance, abstention and policy checking, human verification and benchmark validity, drawn from scholarly databases, publisher and proceedings pages, preprint servers and official agency sites. A source was included when it supported a consequential claim or design choice and could be traced to an opened record, and every cited source is classified by the evidence read: full text or relevant sections, abstract with verified metadata, official document, or news report (used only for recent events and identified as such in the text). Two older sources are retained deliberately: a foundational model of automation and the primary legal source for Timor-Leste's official languages. The search of 22 institutions' AI guidance and of CHED memorandum-order titles is recorded in the project repository. Failure to find a study or a current policy document does not establish that none exists.

Appendix J. Outcome definitions and sensitivity analyses

Scoring. Each criterion is judged Correct, Incorrect, Incomplete or Cannot judge. Equivalent wording and approved alternative readings are accepted. Format failures are not themselves content errors; interpretable raw text is scored on substance. Criterion correctness is the share of a case's criteria judged Correct, averaged over repetitions and then cases, with unread answers counted as zero and a mean over the answers read reported beside it. Over-referral was not recorded as a separate field and is reported only where scoring reasons document it.

Missing and unread answers. Answers that are missing or have not been read stay in the scheduled denominator. Two bounding sensitivities resolve every unresolved answer: least favorable to the checks (withheld unresolved answers treated as acceptable and released ones as serious) and least favorable to direct release (all unresolved answers treated as serious).

Release decisions. Release with a warning counts as release. A routed answer counts as withheld. Serious answers withheld are reported as withheld from ordinary release; handling decisions show whether a person accepted them.

Right-reason inspection. For every serious answer withheld under controls, the check that fired is classified as matched, partly matched or unrelated to the scored error. Blocked acceptable answers are classified by cause (format, quotation not found in the field the check searches, numeric value, unit label, rounding, fraction versus percent, or duplicate entries). The frozen sensitivity based on failure-class matching is not interpreted because failure classes were not collected.

Sensitivity analyses. Each repetition separately; leaving out one institution at a time (a check on case composition, not a comparison of institutions); omitting two cases defined in code as interpretation-sensitive (W1-04, W1-08); omitting the six cases revised after the AI audit; Critical-only severity; excluding all 24 answers seen in the timing exercise, and separately the 12 seen under controlled release; both bounding sensitivities; and the frozen not-scorable-as-unacceptable sensitivity, which cannot raise serious failures and so is not a bound. [Appendix G](#) reports selected results; all are in the supplementary results file, and the text notes any in which a paired count changes direction.

Independent review. Each returned packet is extracted verbatim and compared with the author's judgments on the same answers, which were recorded before any return was opened: exact criterion agreement, agreement on whether an answer has a serious error, and agreement on acceptability (derived for reviewers as severity None or Minor with no undecided criterion), each with its two-by-two counts. Disagreements are reconciled transparently without overwriting either original judgment, and agreement describes only the reviewed subset. One packet arrived with its form controls flattened into ordinary text; its answers were read from the document text in the

same fixed layout, by a reader required to reproduce the form controls exactly on every return that still carried them, and the route used is recorded for each return.

Timing. Primary summaries exclude flagged items ([Appendix H](#)); all recorded times are a sensitivity. Correction time is summarized as the number of items needing correction and the median among them.